

How Much Does Interrupt-Mode `io_uring` Help Shared-Memory IPC?

A Measurement Study of Wakeup Mechanisms for Lock-Free SPSC Rings

Rahul Raman Yuvraj Deshmukh B. Thangaraju

*Department of Computer Science and Engineering
International Institute of Information Technology Bangalore (IIITB)
Bangalore, India*

Email: {rahul.raman, yuvraj.deshmukh, b.thangaraju}@iiitb.ac.in

Abstract—Inter-Process Communication (IPC) performance is a foundational bottleneck in concurrent, multi-process software systems. Traditional kernel-mediated mechanisms—POSIX Pipes, UNIX Domain Sockets, POSIX Message Queues—impose mandatory system calls, context switches, and user-to-kernel memory copying on every message, creating an irreducible per-message overhead of several microseconds. Modern high-performance systems reduce this cost by employing lock-free Single-Producer Single-Consumer (SPSC) shared-memory ring buffers that transfer payload data entirely in userspace via `memcpy` into POSIX shared memory regions. A critical and largely unanswered question, however, is: *when the shared ring empties and the consumer must sleep, which wakeup primitive should be used—and what does Linux’s `io_uring` framework specifically contribute over simpler alternatives such as `futex` or `eventfd`?*

This paper presents a rigorous empirical measurement study that decouples payload transport from wakeup-mechanism signaling and applies a two-suite evaluation framework: (1) a *wakeup-mechanism ablation* comparing six coordination strategies (`busy_poll`, `spin_backoff`, `adaptive`, `futex`, `eventfd`, `io_uring`) on an *identical* cache-aligned SPSC ring under saturated and bursty traffic conditions across a 64 B–1 MiB message-size sweep; and (2) a *corrected depth-1 ping-pong latency suite* in which both the send and receive timestamps are recorded on the single initiator core using `CLOCK_MONOTONIC_RAW`, eliminating cross-core TSC drift and in-flight pipelining artefacts, reporting full percentile distributions (P50, P90, P99, P99.9) over 100 000 round-trips per configuration.

Our results show that the shared-memory ring delivers the primary performance gains. Spinning variants achieve median single-trip latencies of 213–218 ns at 64 B—a 15× improvement over pipe or socket—and peak streaming throughput of 27.88 GiB/s at 64 KiB. In depth-1 ping-pong tests (Suite 2), `io_uring` incurs a modest per-event ring submission overhead (+1.3 μ s over `futex/eventfd` at 64 B). However, under bursty wakeup conditions (Suite 1), where wakeups occur when the ring empties, `io_uring` performs on par with `futex` and `eventfd` (≈ 1.09 – 1.10 ms median wakeup latency with ≈ 0 % idle CPU consumption). `io_uring`’s key architectural advantage lies in providing a unified, future-proof asynchronous I/O interface for multi-channel event loops without sacrificing bursty wakeup efficiency.

Index Terms—`io_uring`, IPC, shared memory, SPSC ring buffer, lock-free, wakeup mechanism, latency, throughput, Linux kernel, systems performance, C++17

I. INTRODUCTION

Modern software infrastructure—spanning high-frequency trading (HFT) matching engines, real-time database kernels, media processing pipelines, and containerised microservice mesh proxies—relies heavily on efficient Inter-Process Communication. Even a modest per-message overhead of a few microseconds accumulates to tens of milliseconds when millions of messages flow per second. At 10 Mop/s, a 3 μ s per-message syscall overhead alone consumes 30 % of the CPU budget.

Linux provides a rich set of mature IPC primitives. POSIX FIFOs and UNIX domain sockets route all data through the kernel, imposing context switches and double memory copies on every message. POSIX message queues add priority ordering but retain kernel-mediated delivery semantics. Shared memory avoids the data-copy penalty by mapping identical physical pages into both producer and consumer address spaces; the kernel is contacted only for initial setup.

The cost of these kernel interactions is non-trivial and grows with message rate. A `write/read` pair on a POSIX pipe costs approximately 1–3 μ s in system-call overhead on modern Linux, independent of payload size. For a system exchanging 5 Mmsg/s, this overhead alone consumes a full CPU core. Beyond raw syscall cost, kernel buffers are allocated from a different NUMA domain relative to the application heap on multi-socket machines, adding DRAM access latency to every message. Real-time systems with sub-1 ms SLAs cannot tolerate this overhead when message rates exceed a few hundred thousand per second.

Lock-free SPSC ring buffers—popularised by the LMAX Disruptor [5] and Intel DPDK [6]—take shared memory a step further by coordinating access using only atomic load/store instructions on cache-line-aligned indices, avoiding any kernel involvement on the data path. Message delivery in the uncontended case reduces to a `memcpy` plus two atomic index updates, costing roughly 200 ns at 64 B—an order of magnitude below any kernel-mediated mechanism.

A remaining challenge is the *empty-ring wakeup problem*:

when the ring is empty, busy-spinning wastes a full CPU core; sleeping the consumer requires a kernel-level wakeup mechanism.

Linux’s `io_uring` interface (introduced in kernel 5.1 [1]) has attracted substantial interest for its promise of zero-syscall I/O. A growing body of practitioner literature presents “`io_uring`-based IPC rings” as offering superior latency and throughput over both traditional IPC and simpler shared-memory approaches. However, these claims conflate two independent design decisions: (1) the payload data path choice (shared memory vs. kernel-copy), and (2) the wakeup mechanism choice (`io_uring` vs. `futex` vs. `spin`). No prior measurement study isolates these dimensions on identical hardware with a controlled ablation.

Contributions

- 1) **Controlled wakeup ablation.** Six pluggable wakeup strategies are implemented on a bit-identical SPSC ring buffer; their impact on latency, throughput, CPU utilization, and syscall rate is measured across a 64 B–1 MiB payload sweep.
- 2) **Corrected depth-1 ping-pong protocol.** Queue depth is enforced at 1 and both round-trip endpoints are timestamped on the same core with `CLOCK_MONOTONIC_RAW`, eliminating TSC skew and pipelining artefacts.
- 3) **Full tail-latency characterisation.** P50, P90, P99, and P99.9 percentiles are reported over 100 000 round-trips per configuration, providing the tail visibility required for SLA-aware system design.
- 4) **Reproducible artifact.** All source code, raw CSV data, and figure-generation scripts are published at https://github.com/Rahul09123/io_uring-IPC.

II. BACKGROUND & RELATED WORK

A. Traditional Kernel-Mediated IPC

POSIX Pipes (FIFOs) (Stevens and Rago [2]) provide uni-directional byte streams backed by a kernel circular buffer (default 64 KiB, tunable to 1 MiB via `fcntl(F_SETPIPE_SZ)`). Every write copies data from user to kernel space; every read copies it back. The producer and consumer alternate between running and sleeping states as the pipe buffer fills and drains, incurring scheduler-wakeup latency on each transition. At very small messages (64 B), the pipe’s fixed per-syscall entry cost dominates; at large messages (≥ 256 KiB), buffer-saturation and block-layer flushing limit throughput to ≈ 5 –6 GiB/s regardless of CPU speed.

UNIX Domain Sockets bypass the network stack but retain the socket buffer abstraction (`sk_buff`). Each `sendmsg/recvmmsg` pair requires a system call and a socket-layer serialisation lock (`mutex_lock` visible in flamegraph traces). Socket send/receive buffers are tuned to 2 MiB via `SO_SNDBUF/SO_RCVBUF` in our implementation. Despite the per-call overhead, UNIX sockets achieve competitive throughput at large payload sizes because the kernel socket

buffer is not bounded by a fixed slot count and can leverage zero-copy page-flipping for contiguous large writes.

POSIX Message Queues (mqqueue) deliver structured, priority-sorted messages via a virtual filesystem inode under `/dev/mqueue`. A key optimisation—`pipelined_send`—delivers directly into a blocked receiver’s address space when one is already waiting, bypassing intermediate queue staging and making `mqqueue` competitive with shared-memory rings at small message sizes. However, the default maximum message size is 8 KiB and the maximum queue depth is 10 messages, both controlled by kernel tunables (`fs.mqueue.msgsize_max`, `fs.mqueue.msg_max`), imposing hard architectural limits at scale.

TABLE I
Kernel-IPC System-Call and Copy Costs per Message

Mechanism	Syscalls /msg	Copies /msg	Kernel buffer
POSIX Pipe	2	2	Ring buf (64 KiB–1 MiB)
UNIX Socket	2	2	<code>sk_buff</code> (up to 2 MiB)
POSIX MQ	2	1–2	VFS inode (10 msgs)
SHM Ring	0	1	None (userspace)

B. Lock-Free SPSC Shared-Memory Rings

Shared-memory IPC maps identical physical pages into producer and consumer address spaces via `shm_open` and `mmap`. SPSC ring buffers (Lamport [4], Herlihy and Shavit [9], LMAX Disruptor [5], DPDK `rte_ring` [6]) coordinate using a pair of atomic integer indices. The critical correctness rule is: the producer advances `head` with `release` ordering after writing payload; the consumer reads `head` with `acquire` ordering before consuming.

The C++17 memory model formally defines these orderings. A *release store* ensures all prior writes are visible to any thread that subsequently performs an *acquire load* of the same variable. For an SPSC ring, this means the consumer observing `head > tail` is guaranteed to see the complete payload written by the producer before `head` was advanced. Using `memory_order_relaxed` on `head` (a common optimisation) works on x86 due to Total Store Order (TSO) but is technically undefined behaviour on weakly-ordered architectures (ARM, RISC-V, Power) that permit Store-Load reordering. We use `memory_order_seq_cst` on the publish store to maintain portability and prevent the lost-wakeup race.

False sharing—where producer and consumer cache lines collide—is prevented by placing `head`, `tail`, and control flags on separate 64-byte-aligned cache lines using `alignas(64)`. Without this separation, a producer store to `head` invalidates the cache line containing `tail` on the consumer core, forcing a cache-coherency RFO (Request For Ownership) on every slot exchange and doubling effective memory latency.

C. The *io_uring* Interface

Introduced by Jens Axboe [1] and merged in Linux 5.1, *io_uring* exposes two ring buffers shared between userspace and the kernel: the Submission Queue (SQ) and the Completion Queue (CQ). In the default interrupt-driven mode (*flags*=0), a single *io_uring_enter* system call submits all pending SQEs and optionally waits for completions. In SQPOLL mode (IORING_SETUP_SQPOLL), a dedicated kernel poller thread eliminates the syscall entirely but requires elevated privileges.

In our implementation, *io_uring* is used *exclusively for the wakeup signaling path*. When the consumer sleeps, it submits an IORING_OP_READ on a named FIFO and blocks in *io_uring_submit_and_wait*. When the producer detects *consumer_sleeping* = 1, it submits an IORING_OP_WRITE to the same FIFO. The payload data path is pure *memcpy* into shared memory—*io_uring* never touches message bytes.

D. Prior Work on *io_uring* Performance

Didona et al. [7] compared *libaio*, *SPDK*, and *io_uring* for storage I/O, demonstrating that polling makes *io_uring* competitive with *SPDK* for sequential workloads. Ren and Trivedi [8] demonstrated that *io_uring* can match kernel-bypass performance for certain NVMe access patterns. Neither study addresses the question of wakeup-primitive selection for shared-memory IPC, leaving this dimension unexplored. Our work fills this gap with a controlled ablation study on a single fixed ring design.

III. SYSTEM DESIGN

Our architecture cleanly separates the data path from the control path.

A. Lock-Free SPSC Ring Buffer Layout

The shared ring buffer is a single POSIX shared memory object opened via *shm_open* and mapped with *mmap* into both the producer and consumer address spaces.

Listing 1. Cache-aligned SPSC RingBuffer (common.h)

```
struct alignas(64) RingBuffer {
    alignas(64) atomic<uint64_t> head;    // Producer write
                                         index
    alignas(64) atomic<uint64_t> tail;    // Consumer read
                                         index
    alignas(64) atomic<uint32_t> consumer_sleeping;

    struct alignas(64) Slot {
        uint64_t send_ns;                // Producer timestamp
        (CLOCK_MONOTONIC_RAW)
        uint64_t wakeup_publish_ns;      // Wakeup-signal timestamp
        uint32_t size;                   // Payload bytes written
        char data[MAX_PAYLOAD];          // Up to 1 MiB of payload
    };
    Slot slots[NUM_SLOTS];               // Fixed 64-slot ring
                                         (~64 MiB total)
};
```

Three design properties are critical.

(1) Cache-line isolation. *head*, *tail*, and *consumer_sleeping* each occupy their own 64-byte cache line. Without this separation, a producer write to *head* invalidates the cache line containing *tail* on the consumer core, forcing

an RFO (Request For Ownership) round-trip on every slot transition. Our hardware counter data (Table X) confirms this: the SHM Ring’s 4.97% LLC miss rate compares favourably to 30–34% for mechanisms that traverse kernel buffers and inadvertently evict ring control variables.

(2) Sequential consistency at publish boundaries. *head* is stored with *memory_order_seq_cst* to prevent Store-Load reordering that could cause the producer to read stale *consumer_sleeping* = 0 and skip sending a wakeup, silently deadlocking the consumer. On x86 (TSO), this is equivalent to a plain store followed by a *mfence*; on ARM64 it compiles to a *stlr* (store-release) instruction.

(3) Lost-wakeup prevention. A classic race exists in any producer-consumer pair: the consumer checks *head* == *tail* (empty), decides to sleep, but before setting *consumer_sleeping*, the producer writes a slot and checks *consumer_sleeping* = 0 (decides not to signal), then the consumer sets *consumer_sleeping* = 1 and waits—forever. We eliminate this race by having the consumer set *consumer_sleeping* = 1 *first*, then re-check *head* != *tail* under a full memory fence. If a slot arrived in the window, the re-check catches it and the consumer proceeds without sleeping.

Data-flow walkthrough. On every message transfer: (a) the producer calls *memcpy* into *slots[head % NUM_SLOTS].data*, writes the timestamp and size, then executes a *seq_cst* store to *head*; (b) the consumer spin-reads *head* with acquire ordering; once it observes *head* > *tail*, it processes the slot and advances *tail* with a release store; (c) if the consumer finds the ring empty after its deadlock-prevention recheck, it invokes the configured wakeup primitive to sleep. Steps (a) and (b) involve zero system calls; step (c) invokes at most one system call per wakeup event.

B. Pluggable Wakeup Layer: Six Ablation Variants

To isolate the contribution of the wakeup mechanism, we implement six interchangeable wakeup modules on a single, bit-identical SPSC ring buffer data structure. Each primitive is invoked via its standard Linux kernel system interface (Linux man-pages [10]): Only the functions *consumer_wait()* and *producer_signal()* differ, guaranteeing that any observed latency difference is attributable solely to the wakeup mechanism.

a) 1. *busy_poll*:. The consumer spins in a tight while(*head.load(relaxed)* == *tail*) loop with no backoff. Zero system calls are issued. Latency is minimal (≈ 200 ns) but one CPU core is permanently pegged at 100%.

b) 2. *spin_backoff*:. Same tight spin with an *_mm_pause()* intrinsic on every iteration. *_mm_pause* inserts a brief pipeline delay on x86, reducing memory-ordering storms and lowering power by $\approx 20\%$ vs. *busy_poll* with negligible latency impact.

c) 3. *adaptive*:. The consumer spins for up to 500 iterations; if no data arrives, it falls back to a *futex* sleep. This hybrid targets low-latency workloads where the ring is

rarely empty (spin wins) while gracefully handling burst-idle traffic (futex sleep saves power).

d) 4. futex.: The consumer stores 1 to `consumer_sleeping` and calls `SYS_futex(FUTEX_WAIT)`. The kernel checks the address atomically; if it still reads 1, the thread is suspended. The producer stores 0 and calls `FUTEX_WAKE`. One system call per sleep and per wakeup event.

e) 5. eventfd.: The consumer polls an eventfd file descriptor and then reads to consume the wakeup token. The producer writes 1 to the eventfd. The eventfd is compatible with `epoll/select`, easing integration into existing event-loop frameworks.

f) 6. io_uring.: The consumer submits an `IORING_OP_READ SQE` targeting a named FIFO (`/tmp/uring_sig_fifo`) and blocks in `io_uring_submit_and_wait(ring, 1)`. The producer submits an `IORING_OP_WRITE SQE` to the same FIFO and calls `io_uring_submit`. The `io_uring` context can simultaneously monitor additional FDs, making it advantageous for frameworks managing many concurrent channels.

TABLE II
Wakeup Mechanism Properties Summary

Variant	Wait Strategy	Idle CPU	Sys/wake
<code>busy_poll</code>	Tight spin	100 %	0
<code>spin_back</code>	Spin+PAUSE	High	0
<code>adaptive</code>	Spin-then-futex	Low	0–1
<code>futex</code>	FUTEX_WAIT	≈0 %	2
<code>eventfd</code>	poll+read	≈0 %	2
<code>io_uring</code>	sub_and_wait	≈0 %	2

IV. EXPERIMENTAL METHODOLOGY

A. Two-Suite Benchmark Framework

Suite 1 — Wakeup-Mechanism Ablation. The producer streams messages continuously under two regimes: *saturated* (back-to-back sends; ring rarely empties) and *bursty* (64-message burst, then 1 ms sleep; ring empties every burst). The bursty regime directly exposes wakeup primitive latency. Payload sizes sweep across {64 B, 256 B, 1 KiB, 4 KiB, 64 KiB, 1 MiB}. $N = 15$ measured runs per configuration (one warmup discarded).

Suite 2 — Corrected Depth-1 Ping-Pong. Queue depth is strictly enforced to 1: the initiator sends one message and blocks waiting for the echo server to reflect it before sending the next. Both t_{start} and t_{end} are recorded on **Initiator Core 1** using `clock_gettime(CLOCK_MONOTONIC_RAW)`. Single-trip latency:

$$L_{\text{trip}} = \frac{t_{\text{end}} - t_{\text{start}}}{2} \quad [\mu\text{s}]$$

100 000 round-trips per configuration; 1 000 warmup trips discarded. P50, P90, P99, and P99.9 percentiles reported.

B. Hardware and System Configuration

TABLE III
Reference Test Environment

Parameter	Value
System	ASUS Vivobook K3605ZF
OS	Ubuntu 24.04.4 LTS (kernel 6.x)
CPU	Intel Core i5-12500H (Alder Lake, 12th Gen)
Architecture	Hybrid P/E-core
Cores/Threads	12 Physical / 16 Logical
Frequency	400–4500 MHz (boost)
L1d Cache	448 KiB (12 instances)
L2 Cache	9 MiB (6 instances)
L3 Cache	18 MiB (shared)
RAM	16 GiB DDR4-3200
Storage	512 GB Samsung NVMe SSD
Compiler	g++ -O2, C++17
liburing	System package, Ubuntu 24.04
CPU Governor	Default (schedutil)

Core affinity. Producer/Initiator pinned to Core 1; Consumer/Echo-Server pinned to Core 2 via `sched_setaffinity`. Both cores share the 18 MiB L3 cache, keeping the ring buffer warm across the inter-process boundary.

CPU tuning and DVFS effects. The reference system runs under standard default Linux conditions (`schedutil` governor, dynamic frequency scaling) as described by Gregg [3]. Consequently, the ≈1.09 ms bursty-regime wakeup latencies partly reflect standard CPU frequency ramp-up from idle C-states alongside raw wakeup primitive overhead. In hard real-time environments, locking CPU frequencies (performance governor, turbo disabled, `isolcpus/nohz_full` core pinning) eliminates DVFS frequency ramp-up delay, isolating pure kernel sleeping and scheduling wakeup primitive latency.

C. Statistical Methodology

For each (mechanism, payload size) pair we compute:

$$\bar{T} = \frac{1}{N} \sum_{k=1}^N T_k \quad [\text{GiB/s}], \quad \bar{L} = \frac{1}{M} \sum_{i=1}^M L_i \quad [\mu\text{s}]$$

$$\sigma_L = \sqrt{\frac{1}{M} \sum_{i=1}^M (L_i - \bar{L})^2}$$

where $N = 15$ is the run count and M is messages per run. 95 % confidence intervals use the Student- t approximation on run-level means.

D. Workload Volumes

TABLE IV
Workload Volumes per (Mechanism, Size) Configuration

Mechanism	≤1 KiB	4–64 KiB	≥256 KiB
POSIX Pipe	32 MiB	256 MiB	2 GiB
POSIX MQ	32 MiB	256 MiB	2 GiB
UNIX Socket	2 GiB	2 GiB	2 GiB
SHM Ring	2 GiB	2 GiB	2 GiB

V. EMPIRICAL RESULTS

A. Unloaded Ping-Pong Latency (Suite 2)

Table V reports single-trip latency distributions (P50, P90, P99, P99.9, and mean $\pm$ 95% CI) at 64 B and 4 KiB, and P50/P99 at 64 KiB and 1 MiB (Section VII). Figures 1 and 2 show the full four-size sweep graphically.

TABLE V

Single-Trip Ping-Pong Latency Distributions (μ s). Full distributions (P90, P99.9, mean $\pm$ 95% CI) at 64 B and 4 KiB; 64 KiB and 1 MiB report P50/P99 only (Threat 7).

Transport / Variant	P50	P90	P99	P99.9	Mean $\pm$ CI
<i>Message Size: 64 B — lower is better</i>					
shm (adaptive)	0.213	0.280	0.582	1.807	0.233 $\pm$ 0.002
shm (busy_poll)	0.218	0.224	0.306	0.675	0.222 $\pm$ 0.001
shm (spin_back)	0.411	0.727	0.781	2.629	0.491 $\pm$ 0.004
unix_socket	3.029	4.049	5.087	7.729	3.237 $\pm$ 0.007
shm (eventfd)	3.113	3.244	4.431	7.576	3.146 $\pm$ 0.002
shm (futex)	3.207	3.409	4.570	7.681	3.252 $\pm$ 0.002
pipe	3.294	3.445	4.669	6.617	3.326 $\pm$ 0.002
posix_mq	3.351	3.581	4.605	6.137	3.400 $\pm$ 0.003
shm (io_uring)	4.542	4.786	6.854	9.086	4.624 $\pm$ 0.004
<i>Message Size: 4 KiB — lower is better</i>					
shm (busy_poll)	1.286	1.363	1.451	3.664	1.299 $\pm$ 0.002
shm (adaptive)	1.286	1.351	1.427	3.646	1.302 $\pm$ 0.008
shm (spin_back)	1.483	1.535	1.733	3.790	1.498 $\pm$ 0.002
shm (futex)	3.746	4.010	5.767	7.807	3.830 $\pm$ 0.003
shm (eventfd)	3.750	4.005	5.273	7.104	3.830 $\pm$ 0.012
pipe	3.922	4.217	6.084	8.015	4.024 $\pm$ 0.004
posix_mq	4.606	4.850	6.171	7.932	4.685 $\pm$ 0.004
unix_socket	5.008	5.491	7.433	10.834	4.867 $\pm$ 0.035
shm (io_uring)	5.369	5.698	7.888	10.660	5.498 $\pm$ 0.005
<i>Message Size: 64 KiB — P50/P99 only (Threat 7)</i>					
shm (busy_poll)	12.995	—	—	—	15.673
shm (adaptive)	13.476	—	—	—	15.657
shm (spin_back)	13.311	—	—	—	15.771
shm (eventfd)	14.831	—	—	—	17.498
shm (futex)	14.914	—	—	—	17.941
unix_socket	15.674	—	—	—	19.092
shm (io_uring)	16.171	—	—	—	19.525
pipe	19.822	—	—	—	23.522
posix_mq	N/A	—	—	—	N/A
<i>Message Size: 1 MiB — P50/P99 only (Threat 7)</i>					
unix_socket	137.911	—	—	—	225.035
shm (spin_back)	149.975	—	—	—	191.538
shm (futex)	152.156	—	—	—	190.628
shm (adaptive)	152.353	—	—	—	180.919
shm (busy_poll)	152.540	—	—	—	214.523
shm (eventfd)	152.591	—	—	—	190.589
shm (io_uring)	154.707	—	—	—	212.985
pipe	261.026	—	—	—	393.990
posix_mq	N/A	—	—	—	N/A

Key observations:

(1) **15 $\times$ spinning advantage at small messages.** At 64 B, `adaptive` (0.213 μ s) and `busy_poll` (0.218 μ s) are 15 $\times$ lower than any kernel-sleeping variant. This gap represents scheduler wakeup latency: transitioning a sleeping thread to running requires dequeuing, CPU selection, and context switching—costing 2–4 μ s on a loaded Linux system.

(2) **Kernel-sleeping variants converge regardless of primitive.** `futex` (3.207 μ s), `eventfd` (3.113 μ s), and `io_uring` (4.542 μ s) all converge to 3–4.5 μ s P50 at 64 B. The slight elevation of `io_uring` reflects SQ/CQ ring entry overhead. (These figures are recomputed from the full 100,000-trip raw




Fig. 1. Depth-1 Ping-Pong P50 Latency vs. Message Size. Spinning SHM variants achieve 213–218 ns at 64 B, 15 $\times$ below kernel-mediated mechanisms. UNIX Socket achieves best latency at 1 MiB (137.9 μ s) due to kernel socket buffer handling large payloads without fixed-slot recycling stalls.

distribution underlying Table V; the `io_uring` value differs marginally from earlier aggregate-based estimates reported in preliminary analysis.)

(3) **POSIX MQ silently fails at large messages.** POSIX MQ returns an error for payloads ≥ 16 KiB due to the kernel default `fs.mqueue.msgsize_max = 8192 B`. These appear as N/A in the table.

(4) **UNIX Socket wins at 1 MiB.** 137.91 μ s vs. ≈ 152 μ s for SHM variants. The kernel socket buffer handles large writes via page-flipping, bypassing the fixed-slot recycling stall that limits the SHM ring.

B. Ablation: Wakeup Variant Comparison (Suite 2)

Figure 3 isolates the six SHM wakeup variants, removing kernel-IPC baselines for clarity. At 1 MiB, all six converge near 150–155 μ s (Table V), confirming memory bandwidth as the common bottleneck.

The 1 MiB convergence is instructive: copying 1 MiB across the ring dominates all wakeup mechanism overhead, making wakeup selection irrelevant at that message size. Wakeup mechanism matters only when message service time is short enough to be latency-dominated.

C. Streaming Throughput: Saturated Regime (Suite 1)

Table VI presents mean throughput; Figure 6 visualises the full size sweep.

TABLE VI

Streaming Throughput (GiB/s) — Saturated Regime, $N = 15$ Runs

Wakeup Variant	64 B	4 KiB	64 KiB	1 MiB
busy_poll	0.635	18.77	27.88	9.67
spin_backoff	0.640	18.75	27.52	9.66
adaptive	0.583	18.89	27.10	9.44
io_uring	0.370	18.45	27.75	9.19
eventfd	0.632	18.46	26.83	9.38
futex	0.637	17.97	26.69	9.35

Key observations:




Fig. 2. Depth-1 Ping-Pong P99 Tail Latency vs. Message Size. `busy_poll` sustains P99 below $0.31\ \mu\text{s}$ at 64 B. All kernel-sleeping variants cluster at $4.4\text{--}6.4\ \mu\text{s}$ at 64 B, confirming the kernel scheduler—not the wakeup API—sets the tail latency floor.




Fig. 3. Ping-Pong Latency Across All Six SHM Wakeup Variants. Spinning variants achieve sub-microsecond latency at small sizes; kernel-sleeping variants converge to $3\text{--}5\ \mu\text{s}$. All six converge near $150\text{--}155\ \mu\text{s}$ at 1 MiB, confirming memory bandwidth as the common bottleneck.

(1) Wakeup mechanism does not affect saturated throughput. Under fully saturated load the ring rarely empties, so the wakeup primitive is never invoked. All six variants produce statistically identical throughput at 4 KiB and 64 KiB.

(2) $5.5\times$ advantage over kernel-mediated IPC. $27.88\ \text{GiB/s}$ (SHM ring) vs. $\approx 5.1\ \text{GiB/s}$ (POSIX pipe) at 64 KiB. This gain originates entirely from eliminating kernel copies and context switches on the data path.

(3) 1 MiB throughput inversion. Throughput drops to $9.2\text{--}9.7\ \text{GiB/s}$ because the fixed 64-slot ring stalls the producer

once all slots are occupied (discussed further in Section VI).

(4) Small-message `io_uring` overhead at 64 B. At 64 B payload, `io_uring`'s saturated throughput ($0.370\ \text{GiB/s}$) is lower than `futex/eventfd` ($0.637\ \text{GiB/s}$) and spinning ($0.667\ \text{GiB/s}$). This occurs because per-message SQE/CQE submission and completion processing introduces a fixed per-message overhead that is noticeable when message payloads are small. At 4 KiB and above, this overhead is fully amortized and all variants converge to identical streaming throughput ($\approx 18.5\ \text{GiB/s}$ at 4 KiB and $27.88\ \text{GiB/s}$ at 64 KiB).

D. Bursty Regime: Pure Wakeup Latency Ablation (Suite 1)

In the bursty regime, the consumer empties the ring and enters a sleep state before each burst. Table VII isolates wakeup latency per variant.

TABLE VII
Bursty-Regime Wakeup and End-to-End Latencies at 64 B

Variant	E2E P50	Wakeup P50	Idle CPU	Syscalls
busy_poll	0.13 μ s	0.012 μ s	100 %	0
spin_backoff	0.13 μ s	0.052 μ s	High	0
adaptive	135.8 μ s	1070.5 μ s	Low	0.01
futex	166.7 μ s	1090.5 μ s	≈ 0 %	0.02
eventfd	166.4 μ s	1093.9 μ s	≈ 0 %	0.02
io_uring	151.0 μ s	1104.5 μ s	≈ 0 %	0.02

TABLE VIII
CPU Resource Efficiency per 64 B Message (Bursty Regime)

Variant	Cycles / Msg	Ctx-Sw / Msg	Insn / Msg
futex	9,310	0.0263	8,512
eventfd	9,653	0.0280	8,711
io_uring	11,338	0.0287	10,928
adaptive	21,620	0.0243	12,323
busy_poll	114,789	0.0003	216,267
spin_backoff	202,011	0.0003	158,135

Table VIII quantifies the CPU cost side of the latency/CPU trade-off already visible qualitatively in Table VII. The kernel-sleeping variants (futex, eventfd, iouring) cost 9,300–11,300 cycles/message—roughly an order of magnitude less than the spinning variants (114,789–202,011 cycles/msg), because a sleeping thread consumes no cycles while blocked, whereas a spinning thread burns cycles continuously regardless of whether new data has arrived.

Note that spinbackoff’s higher cycle count relative to busypoll (202,011 vs. 114,789) does not contradict its power-saving design intent from Section ??: the `_mm_pause` intrinsic inserted on every iteration deliberately reduces instructions retired per cycle (158,135 vs. 216,267 insn/msg, i.e. IPC ≈ 0.78 vs. ≈ 1.88), which is what actually lowers power draw at the microarchitectural level—cycle count and power are not the same axis. adaptive sits at an intermediate 21,620 cycles/msg, reflecting its spin-then-block hybrid: most cycles come from the initial spin window, with the fallback futex wait contributing comparatively little once the ring is confirmed empty.

The central ablation finding: io_uring’s wakeup latency (1104.5 μ s P50) is indistinguishable from futex (1090.5 μ s) and eventfd (1093.9 μ s). All three follow the same kernel path: store a flag, issue a syscall, suspend in the scheduler, receive a wakeup interrupt, return to userspace. The 1 ms range is dominated by CPU clock ramp-up from an idle P-state, not by the wakeup API.

TABLE IX

Offered-Load Sweep: Median Wakeup Latency (μ s) at 64 B. Spinning variants maintain <0.1 μ s regardless of load. Kernel-sleeping variants scale from ≈ 1090 μ s at 25% to ≈ 107 μ s at 90%.

Variant	25%	50%	75%	90%	Sat.
busy_poll	0.01	0.01	0.01	0.01	0.01
spin_backoff	0.05	0.05	0.05	0.05	0.05
adaptive	1255.2	524.6	181.6	69.5	0.05
eventfd	1430.7	670.3	258.2	107.6	0.15
futex	1433.0	669.5	257.9	106.5	0.26
io_uring	1452.5	672.8	260.1	107.3	0.77

As shown in Table IX and visualised in Figure 5, under offered-load sweeps (25%, 50%, 75%, 90% of saturated capacity), spinning mechanisms (busypoll and spinbackoff) maintain fixed sub-microsecond latencies regardless of offered load. For kernel-sleeping variants (futex, eventfd, io_uring), median latency dynamically scales down as load increases—dropping from ≈ 1430 – 1452 μ s at 25% load down to ≈ 106 – 107 μ s at 90% load as CPU cores remain in active C-states for longer stretches between inter-arrival bursts. At saturation, io_uring’s slightly higher latency (0.77 μ s) compared to futex/eventfd reflects the constant overhead of SQE/CQE ring management on the non-blocking fast path. adaptive hybrid spinning captures the best of both worlds, reducing low-load sleep penalties while achieving pure spinning performance (0.05 μ s) under saturation.

The bursty-regime protocol (Table VII: 64-message burst, then a fixed 1 ms idle gap) and the offered-load sweep (Table IX: continuous Poisson-distributed arrivals at a fixed fraction of saturated rate) are deliberately different idle-gap structures and are not directly comparable message-for-message. The fixed 1 ms gap in the bursty protocol is shorter than the mean inter-arrival gap implied by 25% of saturated throughput at 64 B, so the CPU has less time to fall to a deeper idle C-state before the next burst arrives—producing the lower ≈ 1090 – 1104 μ s wakeup latency in Table VII relative to Table IX’s 25%-load row. As offered load rises toward saturation in Table IX, the mean inter-arrival gap shrinks below the bursty protocol’s fixed 1 ms gap, which is why the 90%-load row (≈ 106 – 107 μ s) falls below the bursty-regime figures despite both nominally representing “low load” conditions in isolation. The two tables should be read as separate experimental protocols probing the same underlying C-state ramp-up effect from different angles, not as repeated measurements of the same condition.

E. IPC Throughput vs. Kernel-Mediated Baselines

To contextualise the performance of our lock-free SPSC shared-memory ring, we benchmarked it against standard Linux kernel-mediated IPC mechanisms across a payload size sweep from 64 B to 1 MiB under saturated conditions.

Figure 6 compares overall streaming throughput. The SHM ring achieves peak throughput of 27.88 GiB/s at 64 KiB, outperforming POSIX Pipe (5.1 GiB/s) by 5.46 $\times$ and UNIX Domain Socket (4.3 GiB/s) by 6.46 $\times$. At 1 MiB, UNIX Sockets (11.1 GiB/s) and POSIX MQ (10.9 GiB/s) slightly exceed the




Fig. 4. Bursty-Regime Wakeup Latency Across Six SHM Variants. Spinning variants maintain near-zero wakeup latency at 100 % CPU cost. Kernel-sleeping variants (*futex*, *eventfd*, *io_uring*) converge to $\approx 1090\text{--}1105\ \mu\text{s}$ under normal conditions (dominated by CPU frequency ramp-up from idle states).




Fig. 5. Throughput by Variant and Arrival Regime. Throughput collapses from $\approx 14\text{ GiB/s}$ under saturated load to near-zero under offered-load sweeps as the ring empties between arrivals, regardless of wakeup mechanism.

SHM ring (9.7 GiB/s) due to fixed ring-slot recycling stalls (analysed in Section VI).

Figure 7 shows median (P50) queued latency on a logarithmic scale under saturation. Traditional kernel-mediated mechanisms exhibit severe latency degradation at small payload sizes due to accumulated system-call overhead, context switches, and queue staging. POSIX Pipe reaches $5,102\ \mu\text{s}$ median queued latency at 64 B under saturation, whereas the SHM ring maintains sub-microsecond latency.

Figure 8 presents the throughput speed-up ratio of the SHM ring relative to each kernel baseline. The SHM ring delivers maximum speedup at 256 B ($6.46\times$ over UNIX Sockets). Speedup remains above $4.5\times$ across 64 B–64 KiB payloads before dipping to $0.87\times$ at 1 MiB due to the 64-slot ring-depth

constraint.

F. Cache and Memory Hierarchy Analysis

TABLE X
LLC Miss Rates Under Bursty Regime at 64 B

Mechanism	L1 Miss %	LLC Miss %	Relative to Pipe
Pipe	2.42 %	33.6 %	1.00 $\times$
Socket	1.12 %	31.5 %	0.94 $\times$
MQ	N/A	3.2 %	0.10 $\times$
SHM Ring	3.97 %	4.97 %	0.15 $\times$

The 4.97 % LLC miss rate confirms that the 64-slot ring working set fits within the 18 MiB L3 cache. Pipe and




Fig. 6. Throughput (GiB/s) vs. Message Size. SHM Ring peaks at 27.88 GiB/s at 64 KiB vs. Pipe’s 5.1 GiB/s. At 1 MiB, UNIX Socket (11.1 GiB/s) and POSIX MQ (10.9 GiB/s) overtake the ring (9.7 GiB/s) due to fixed ring-slot depth.




Fig. 7. Median (P50) Latency (μ s, log scale) vs. Message Size under saturated load. These are queued-latency measurements (not corrected unloaded latency). POSIX Pipe reaches 5,102 μ s at 64 B due to accumulated queuing and context-switch overhead.




Fig. 8. Throughput Speed-Up Ratio (SHM Ring $\div$ Baseline). Peak speedup is $6.46\times$ vs. UNIX Socket at 256 B. Below parity at 1 MiB ($0.87\times$ vs. Socket) due to the fixed 64-slot ring-depth constraint.

Socket push data through kernel buffers evicted from L3 between accesses, leading to 30–34 % LLC miss rates. The alignas(64) cache-line isolation ensures a producer store to head never invalidates the consumer’s tail cache line.




Fig. 9. L1 D-Cache and LLC Miss Rates by IPC Mechanism. Pipe and Socket exhibit $\approx 30\text{--}34\%$ LLC miss rates due to kernel buffer staging evicting ring cache lines. SHM Ring (4.97 %) and POSIX MQ (3.2 %) achieve low LLC miss rates by keeping data in the shared 18 MiB L3.

VI. DISCUSSION

A. What *io_uring* Actually Contributes to IPC

The ablation experiment answers the title question directly: on a single SPSC ring channel, ***io_uring* wakeup provides no latency advantage over *futex* or *eventfd***. In strict depth-1 ping-pong (Suite 2), *io_uring* produces 4.54 μ s P50 at 64 B (+1.3 μ s over *futex*’s 3.21 μ s) due to per-event SQE/CQE submission overhead. However, under bursty wakeup conditions (Suite 1), where wakeups occur when the ring empties, all three kernel-sleeping primitives yield statistically indistinguishable median wakeup latency ($\approx 1.09\text{--}1.10$ ms) and $\approx 0\%$ idle CPU overhead.

The primary performance advantage attributed to “*io_uring* IPC” in practitioner benchmarks originates entirely from replacing kernel-copy data paths with shared-memory *memcpy*. When the data path is held constant (as in our ablation), the wakeup mechanism accounts for $<0.1\%$ of total transfer time under saturated load and contributes only to per-burst overhead under bursty load. Table VI confirms this: all six variants achieve 26.7–27.9 GiB/s at 64 KiB under saturated load, a spread of less than 5 % attributable to measurement noise rather than wakeup semantics.

A useful mental model is to view SPSC IPC performance as a two-layer stack:

- 1) **Data path layer** (dominant): shared memory vs. kernel-copy. This layer determines 95 %+ of throughput and latency at message sizes where payload transfer dominates.
- 2) **Wakeup layer** (secondary): spinning vs. blocking. This layer dominates only when message rate falls below the spin-wait threshold—i.e., when the ring is genuinely empty between bursts.

Designing “*io_uring* IPC” as if it were a single unified mechanism misrepresents this two-layer structure. The correct framing is: shared-memory IPC with *io_uring* wakeup signaling. The first clause explains the speedup; the second clause explains the engineering interface choice.

io_uring’s genuine architectural advantages over *futex*/*eventfd* for this use case are:

- 1) **Unified multi-FD event loop.** A single `io_uring` context concurrently monitors many FDs without `epoll` overhead. For a microservice managing 100 IPC ring channels, one `io_uring_submit_and_wait` replaces 100 individual `futex_wait` calls, reducing per-iteration syscall count from $O(N)$ to $O(1)$.
- 2) **Batched submission.** Multiple wakeup signals destined for different channels can be coalesced into a single `io_uring_enter`, amortizing kernel-entry cost across B signals. At batch size $B = 16$, the effective per-signal syscall cost drops to $1/16^{\text{th}}$ of a standalone `FUTEX_WAKE`—a meaningful saving at $>1\text{Mwakeup/s}$.
- 3) **Regime-dependent latency trade-off.** In strict depth-1 ping-pong scenarios (Suite 2), `io_uring` incurs a $+1.3\mu\text{s}$ per-event SQE/CQE ring management overhead compared to `futex/eventfd`. However, in bursty or multi-channel event loops (Suite 1), where wakeups occur when the ring empties, `io_uring` matches `futex` latency ($\approx 1.09\text{ms}$) while scaling efficiently to many concurrent channels.
- 4) **SQPOLL upgrade path.** Enabling `IORING_SETUP_SQPOLL` requires a single flag change and eliminates the `io_uring_enter` syscall entirely for the wakeup path, placing `io_uring` between spinning and sleeping variants in the latency spectrum without dedicated-core overhead.

B. Architectural Guidance by Use Case

a) *High-Frequency Trading / Sub- μs Latency.*: Use SHM ring with `busy_poll` or `adaptive`. Both achieve $P50 \approx 213\text{ns}$ and $P99 \leq 0.6\mu\text{s}$ at 64 B. The 100% idle CPU cost is acceptable when a dedicated core per channel is standard practice. `adaptive` is preferred for intermittent workloads: it spins for up to 500 iterations (covering sub-200ns wakeup cases) and falls back to `futex` only when the ring is genuinely empty, reducing power consumption by up to 80% during idle gaps without hot-path latency impact.

b) *Cloud Microservices / Energy-Constrained Environments.*: Use SHM ring with `futex` or `eventfd`. Both deliver $\approx 0\%$ idle CPU, sub- $5\mu\text{s}$ P50, and sub- $6\mu\text{s}$ P99 at 64 B under normal conditions. `eventfd` integrates naturally into existing `poll/epoll` event loops without `futex`-specific code, making it the lower-friction choice for retrofit into existing stacks.

c) *Async Event-Driven Frameworks.*: Use SHM ring with `io_uring`. Wakeup latency matches `futex/eventfd` precisely ($\approx 3\text{--}4.5\mu\text{s}$ P50), but the `io_uring` context unifies IPC ring monitoring with network I/O, disk I/O, and timer events in a single `submit_and_wait` call. This eliminates the complexity of maintaining parallel `epoll` and `futex` contexts and provides a clear upgrade path to SQPOLL mode.

d) *Large Payloads ($>1\text{MiB}$) or Deep Buffering.*: Use UNIX Domain Sockets. At 1 MiB, UNIX Socket achieves $137.9\mu\text{s}$ P50 vs. $\approx 152\mu\text{s}$ for all SHM ring variants, because the kernel socket buffer uses page-flipping without per-slot recycling stalls.

C. Ring-Depth Constraint at 1 MiB

The throughput inversion at 1 MiB—where UNIX Socket (11.1 GiB/s) and POSIX MQ (10.9 GiB/s) overtake the SHM ring (9.7 GiB/s)—is a direct consequence of the fixed 64-slot ring geometry. With each slot sized at 1 MiB, the ring holds at most 64 MiB of in-flight data. The producer stalls as soon as all 64 slots are occupied while the consumer processes each 1 MiB slot sequentially.

Quantitatively, assume the consumer can drain one 1 MiB slot in time T_c and the producer can fill one slot in time $T_p < T_c$ (producer is faster). Steady-state throughput is bounded by the consumer’s per-slot service rate:

$$\text{Throughput} \leq \frac{1\text{ MiB}}{T_c}$$

When T_c is dominated by the consumer checksum+timestamp operation ($\approx 102.4\mu\text{s}$ per 1 MiB slot at 10 GB/s memory bandwidth), maximum ring throughput is bounded by $\approx 9.7\text{GiB/s}$ —matching our measured value. The fixed $N = 64$ slot count bounds total in-flight buffering ($64 \times 1\text{ MiB} = 64\text{ MiB}$), forcing the producer to stall whenever consumer processing lags behind the producer fill rate. The socket kernel buffer, by contrast, effectively provides unlimited in-flight depth by spilling to physical DRAM, bypassing the slot-count constraint.

Increasing `NUM_SLOTS` to 256 would raise the in-flight capacity to 256 MiB and likely recover the throughput advantage at 1 MiB. The total ring footprint grows from $\approx 64\text{ MiB}$ to $\approx 256\text{ MiB}$ —still viable since producer and consumer access only the most recent few slots in steady state, and the active working set ($\sim 1\text{--}4$ slots) fits comfortably in the 18 MiB L3 cache. This adjustment is a one-line change to `common.h` and does not require any modification to the wakeup layer.

D. Comparison with Related Measurement Studies

Our findings agree with Didona et al. [7] that `io_uring` provides no inherent advantage for unbatched single-operation IPC when compared to simpler synchronous primitives. Similarly, in our depth-1 IPC ping-pong benchmark (Suite 2), per-message submission overhead makes `io_uring` slightly slower than `futex`, matching its latency only when wakeups are amortized across bursty arrivals (Suite 1). In contrast to Ren and Trivedi [8], who attribute speedups to `io_uring`’s reduced per-operation overhead for NVMe I/O, our experiment shows this benefit does not transfer to shared-memory IPC wakeup: the NVMe path benefit arises from batching DMA descriptor submissions, whereas SPSC wakeup is inherently one-at-a-time (one signal per empty-ring event). Batching is only applicable when multiple channels empty simultaneously, which requires a multi-ring architecture outside the scope of this study.

Practitioner implementations such as Aeron [5] use a dedicated “media driver” thread that busy-spins on all channels, effectively implementing `busy_poll` at the framework level. Our data support this design choice: `busy_poll` achieves $0.012\mu\text{s}$ wakeup latency vs. $1090\text{--}1105\mu\text{s}$ for all kernel-sleeping variants. The trade-off is one dedicated core per media

driver; for applications that can afford this, our results confirm it is the optimal strategy.

VII. THREATS TO VALIDITY

- 1) **Single node and microarchitecture.** All experiments ran on one Intel 12th-Gen x86_64 laptop. Results on ARM64, AMD Zen, or multi-socket NUMA may differ due to different cache-coherency protocols and scheduler implementations.
- 2) **SPSC topology only.** Results apply to Single-Producer Single-Consumer queues. MPMC rings require CAS loops introducing contention absent in SPSC.
- 3) **Synthetic payload workload.** Payloads consist of strided byte patterns. Real application payloads (protobuf, encrypted frames) may shift the bottleneck to serialisation rather than ring mechanics.
- 4) **CPU frequency effects.** The ≈ 1.09 ms bursty-regime wakeup latencies reflect standard CPU frequency ramp-up under normal default conditions (`schedutil`). Practitioners running strict hard-real-time systems would typically lock CPU frequencies to maximum to eliminate this ramp-up delay, reducing wakeup latency further.
- 5) **POSIX MQ size cap.** Results for POSIX MQ at ≥ 16 KiB are unavailable due to the default `fs.mqueue.msgsize_max = 8192 B` kernel limit.
- 6) **Interrupt-mode `io_uring` scope.** All `io_uring` experiments evaluate standard default interrupt mode (`flags = 0`). SQPOLL mode (`IORING_SETUP_SQPOLL`) would eliminate the wakeup syscall entirely but requires `CAP_SYS_ADMIN` privileges. Our empirical conclusions are strictly scoped to interrupt-mode `io_uring` wakeup signaling.
- 7) **Full tail-percentile capture range.** Full tail-percentile capture (P50, P90, P99, P99.9) is reported for 64 B and 4 KiB payloads. For 64 KiB and 1 MiB payloads, total round-trip latency is dominated by memory-copy bandwidth ($T_{\text{copy}} \gg T_{\text{wakeup}}$), where tail jitter is driven by DRAM bus contention rather than OS scheduling transitions.

VIII. CONCLUSION & FUTURE WORK

This paper presented a rigorous, controlled measurement study of wakeup mechanism performance in lock-free shared-memory SPSC IPC. By holding the ring buffer design constant and varying only the sleep/wakeup primitive, we isolated the true contribution of `io_uring` to IPC latency.

Summary of findings:

- 1) **The shared-memory data path drives the performance gains.** Spinning on a cache-aligned ring delivers 213 ns P50 at 64 B and 27.88 GiB/s peak throughput—a $15\times$ latency and $5.5\times$ throughput improvement over kernel-mediated IPC, independent of the wakeup primitive installed.
- 2) **`io_uring` matches `futex` and `eventfd` under bursty wakeup conditions.** Under bursty load (Suite 1), `io_uring` matches `futex` and `eventfd` (≈ 1.09 – 1.10 ms wakeup latency), though paying a modest $+1.3\mu\text{s}$ per-event SQE submission overhead in strict depth-1 ping-pong (Suite 2). `io_uring`’s architectural advantages—unified multi-FD event loops, batched submission, SQPOLL upgrade path—justify its adoption without sacrificing bursty wakeup efficiency.
- 3) **Spinning variants dominate sub- μs latency requirements.** `adaptive` is the preferred choice: it achieves 213 ns hot-path latency and gracefully reduces CPU consumption during idle periods.
- 4) **The 1 MiB throughput inversion is a ring-depth artefact.** Increasing `NUM_SLOTS` from 64 to 256 is expected to recover the throughput advantage at 1 MiB with minimal footprint increase.
- 5) **Cache hierarchy alignment is critical.** The SHM ring achieves a 4.97 % LLC miss rate vs. 30–34 % for kernel-copy mechanisms, confirming that `alignas(64)` false-sharing prevention is essential to realising the shared-memory advantage.

Future directions:

- **SQPOLL mode.** Enabling `IORING_SETUP_SQPOLL` would eliminate the `io_uring_enter` syscall, placing it between spinning and sleeping variants latency-wise.
- **MPMC extension.** A lock-free MPMC ring with sequence-lock or CAS-based reservation extends the ablation to multi-producer settings.
- **Cross-NUMA evaluation.** Profiling with producer on NUMA node 0 and consumer on NUMA node 1 would quantify LLC miss penalties and motivate NUMA-aware ring partitioning.
- **RDMA baseline.** An InfiniBand one-sided write baseline contextualises these results within HPC cluster IPC.
- **Production traffic models.** Replacing synthetic payloads with real-world Kafka message logs or gRPC frame traces would validate whether serialisation overhead shifts the bottleneck off ring mechanics.

REFERENCES

- [1] J. Axboe, “Efficient I/O with `io_uring`,” Linux Kernel Documentation, Tech. Rep., 2019. Available: https://kernel.dk/io_uring.pdf
- [2] W. R. Stevens and S. A. Rago, *Advanced Programming in the UNIX Environment*, 3rd ed. Addison-Wesley Professional, 2013.
- [3] B. Gregg, *Systems Performance: Enterprise and the Cloud*, 2nd ed. Addison-Wesley, 2020.
- [4] L. Lamport, “Specifying concurrent program modules,” *ACM Trans. Programming Languages and Systems*, vol. 5, no. 2, pp. 190–222, Apr. 1983.
- [5] M. Thompson, D. Farley, M. Barker, P. Gee, and A. Stewart, “Disruptor: High performance alternative to bounded queues for exchanging data between concurrent threads,” LMAX Exchange, White Paper, 2011.
- [6] DPDK Project, “DPDK Programmer’s Guide: Ring Library,” 2024. Available: https://doc.dpdk.org/guides/prog_guide/ring_lib.html
- [7] D. Didona, J. Pfefferle, N. Ioannou, B. Metzler, and A. Trivedi, “Understanding modern storage APIs: A systematic study of `libaio`, `SPDK`, and `io_uring`,” in *Proc. 15th ACM International Systems and Storage Conference (SYSTOR ’22)*, Haifa, Israel, Jun. 2022, pp. 1–16.
- [8] Z. Ren and A. Trivedi, “Performance Characterization of Modern Storage Stacks: `POSIX I/O`, `libaio`, `SPDK`, and `io_uring`,” in *Proc. 3rd Workshop on Challenges and Opportunities of Efficient and Performant Storage Systems (CHEOPS ’23)*, Rome, Italy, May 2023.

- [9] M. Herlihy and N. Shavit, *The Art of Multiprocessor Programming*. Morgan Kaufmann, 2008.
- [10] Linux man-pages project, `futex(2)`, `eventfd(2)`, `io_uring(7)`, `mq_overview(7)`, `pipe(7)`, `unix(7)`, kernel.org, 2024. Available: <https://man7.org/linux/man-pages/>

APPENDIX

RING BUFFER MEMORY LAYOUT

```
// Shared memory: /dev/shm/ipc_ablation_ring (~64 MiB
total)
struct alignas(64) RingBuffer {
    alignas(64) atomic<uint64_t> head;           // cache
        line 0
    alignas(64) atomic<uint64_t> tail;           // cache
        line 1
    alignas(64) atomic<uint32_t> consumer_sleeping; //
        cache line 2

    struct alignas(64) Slot {
        uint64_t send_ns;                       // CLOCK_MONOTONIC_RAW
        producer stamp
        uint64_t wakeup_publish_ns;             // wakeup-signal
        timestamp
        uint32_t size;                           // payload byte count
        char data[MAX_PAYLOAD];                 // up to 1 MiB
    };
    Slot slots[NUM_SLOTS];                     // 64 slots by default
};
// Total footprint: 64 * (64 + 1 MiB) ~= 64 MiB
```

SOURCE CODE REPOSITORY

All source code, benchmark scripts, raw CSV data, and figure-generation scripts for this study are publicly available at:

https://github.com/Rahul09123/io_uring-IPC